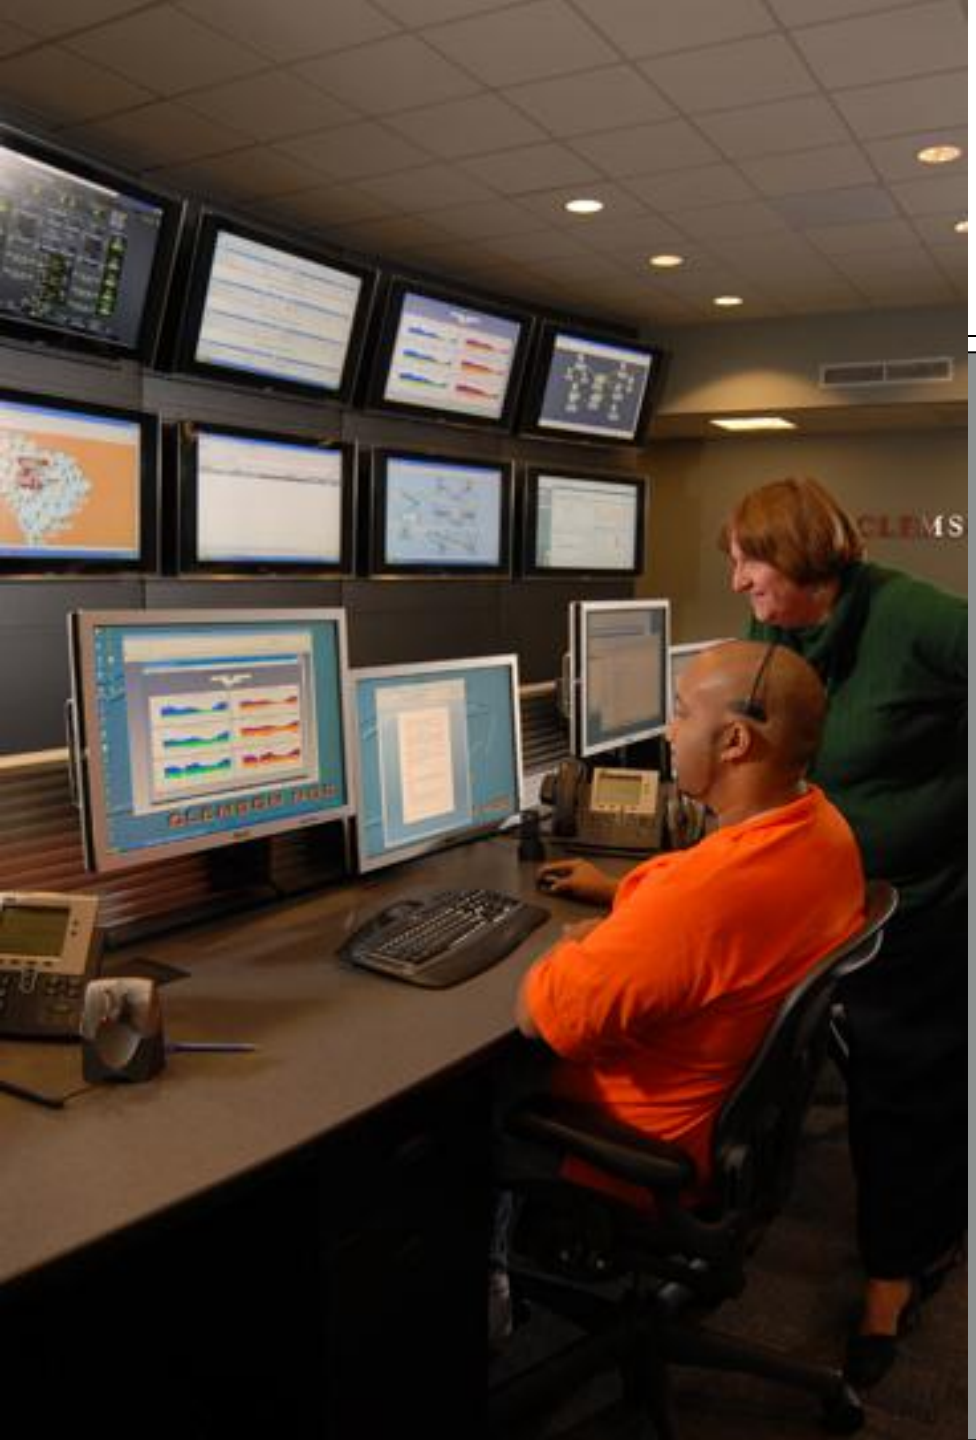

CLEMSON
COMPUTING AND INFORMATION TECHNOLOGY

Introduction to Palmetto

Galen Collier
Computational Scientist

Cyberinfrastructure Technology Integration (CITI)
Clemson University

Jan. 17, 2012

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

The Palmetto Cluster

- Both Shared and Distributed configs
- Operates at over 96 TF/s
- #128 on November 2011 “Top 500” list
- #2 among public academic institutions
- 1,616 compute nodes (14,168 cores)
- Operating System: Scientific Linux 6
- Myrinet 10G network interconnect
- Data Storage:

  115 TB “scratch” filesystem

  > 72 TB of purchased backed-up storage

CLEMSON
COMPUTING AND INFORMATION TECHNOLOGY

Node Specifications (1 of 2)

	Name	Count	Model	Processor	L2 Cache	Cores	Memory	Local Disk
compute node phase 1	node0001-0257	257	Dell PE 1950	Intel Xeon E5345 @2.33GHz x 2	4 MB	8	12 GB	80 GB (SATA)
compute node phase 2	node0258-0515	258	Dell PE 1950	Intel Xeon E5410 @2.33GHz x 2	6 MB	8	12 GB	80 GB (SATA)
compute node phase 3	node0516-0771	256	Sun X2200 M2 x64	AMD Opteron 2356 @ 2.3GHz x 2	4 MB	8	16 GB	250 GB (SATA)
compute node phase 4	node0772-1023,1108-1111	256	IBM dx340	Intel Xeon E5410 @2.33GHz x 2	6MB	8	16 GB	160 GB (SATA)
compute node phase 4.1	node1024-1107	84	IBM dx340	Intel Xeon E5410 @2.33GHz x 2	6MB	8	16 GB	160 GB (SATA)
compute node (former CCMS nodes)	node1112-1541	430	Sun X6250	Intel Xeon L5420 @2.5GHz x 2	6MB	8	32 GB	160 GB (SATA)
compute node phase 6	nodes 1553-1622	70	HP DL 165 G7	AMD Opteron 6172 @2.1GHz x 2	12MB	24	48 GB	250 GB (SATA)

CLEMSON
COMPUTING AND INFORMATION TECHNOLOGY

Node Specifications (2 of 2)

	Name	Count	Model	Processor	L2 Cache	Cores	Memory	Local Disk
regular large shared memory systems	nodelm01-nodelm04	4	HP DL 580 G7	Intel Xeon 7542 @ 2.66 GHz x 4	18MB	24	512 GB	146 GB (SAS)
math sciences large memory	nodemath	1	HP DL 980 G7	Intel Xeon 7560 @ 2.66 GHz x 8		64	2 TB	

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

/scratch

- Temporary “work” directory for all jobs
  (Move all files in, run jobs, move all files out)
- Open to all users, create a directory for yourself
- Specialized high-performance hardware
- Storage server software = OrangeFS
- Designed to handle very high I/O activity
- 115 TB of space, open to all users

Note: Try to keep the number of files per directory to less than 1,000

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Accessing Palmetto

Command-line interface, any Secure Shell (ssh) client

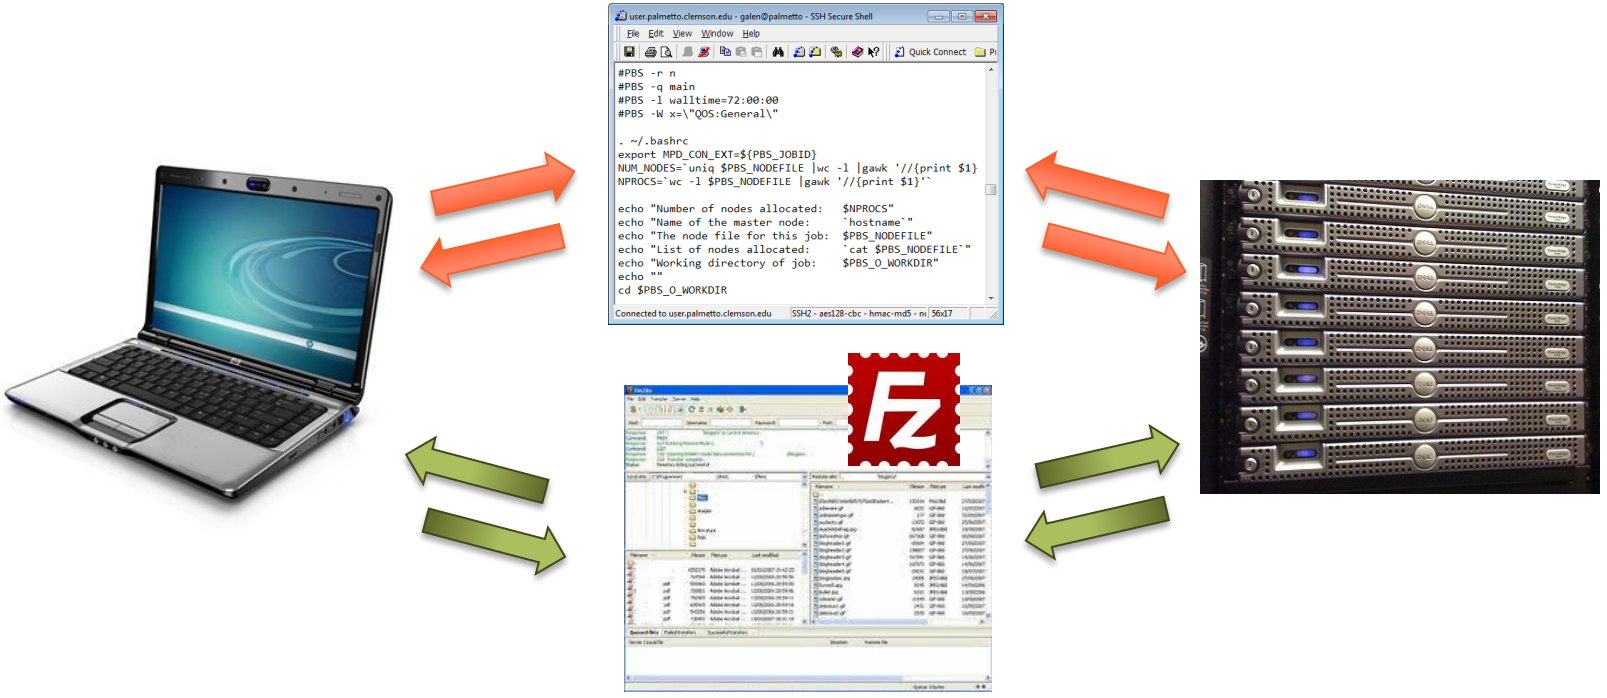

Transfer files to/from using FileZilla (or scp)

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

MOTD (/etc/motd)

```
--------------------------------------------------------------------------
        Welcome to the PALMETTO CLUSTER at CLEMSON UNIVERSITY

  * Please email ithelp@clemson.edu with questions or to report problems.
  * HPC webpage is http://citi.clemson.edu/hpc

  * The FIRST TUESDAY OF EACH MONTH, from 9:00am to 12:00noon,
    is reserved for system wide cluster maintenance.

  DO NOT RUN JOBS ON THE USER LOGIN NODE.  THEY WILL BE
  TERMINATED WITHOUT NOTICE.  NO EXCEPTIONS.

  Useful commands:
    checkquota               - get your current disk quota
    module avail             - list software packages
    qstat -xf jobid          - check status of jobid
    qstat -Qf queuename      - check status of queuename
    pbstop -m 24             - check status of jobs across the cluster
    cat /etc/hardware-table  - list hardware types: ram,cores,chip,etc.

  Palmetto User Guide:   http://desktop2petascale.org/resources/159

  Workaround for PBS and -k option:
  If you use "-k oe" or "-k e" or "-k o", you must set permissions
  on your /home dir to 711:    chmod 711 /home/userid
  otherwise your output will be blank at the job's end.

---------------------------- /etc/motd ----- Last Updated: 10-Jan-2012 ---
```

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Passwordless SSH

```
ssh-keygen -t rsa
```

Generating public/private rsa key pair.
Enter file in which to save the key (/home/userid/.ssh/id_rsa): *[Enter]*
Enter passphrase (empty for no passphrase): *[Enter]*
Enter same passphrase again: *[Enter]*
Your identification has been saved in /home/userid/.ssh/id_rsa.
Your public key has been saved in /home/userid/.ssh/id_rsa.pub.
The key fingerprint is: 64:72:2a:7b:20:fa:a7:0c:91:26:a6:43:85:0b:1c:21

```
cd .ssh
```

```
cp id_rsa.pub authorized_keys
```

(you can also test your configuration: **`qsub -I`**)

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Copy Example Files

Copy the examples to your home (or scratch) directory:

```
cd ~
cp -r /scratch/galen/intro.palmetto ~
```

```
[galen@user001 ~]$ cd /scratch/galen/intro.palmetto
[galen@user001 intro.palmetto]$ ls
```

```
bashrc.example                  gethostname.mpi.c
job.gethostname.bash            job.template.bash
examples                        hello_world_mpi.c
job.gethostname.mpi.bash        lammps-30Sep11
gethostname.c                   hello_world_openmp.c
job.hello_world.omp.bash        pbsdsh
```

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Try an Example Job

- LAMMPS is a molecular dynamics simulation program
- Cooling of a binary mixture
- 5,000 atoms (LJ interactions)
- 50,000 0.005 ps steps

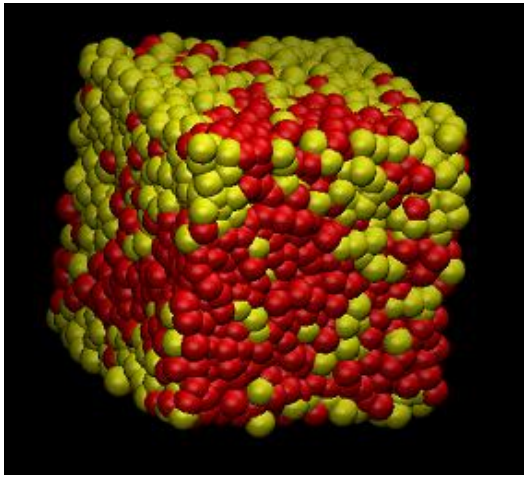

```
cd ~/intro.palmetto/examples/lammps.example
```

Edit job.lammps.bash (optional), then run the job:

```
qsub job.lammps.bash
```

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Or try this one…

- NAMD is a molecular dynamics simulation program
- apoa1 system (over 92,000 atoms)
- 12 Å cutoff, PME every 4 steps

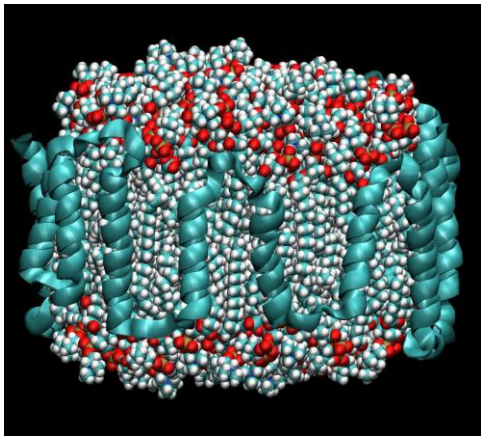

```
cd ~/intro.palmetto/examples/namd.example
```

Edit job.namd.bash (optional), then run the job:

```
qsub job.namd.bash
```

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Setting Your Environment

- To get a list of available software modules:

  `module avail`

- To use specific software every time you log-in, place the **`module add`** command in your ~/.bashrc file:

  `module add intel/12.0 mpich2/1.4`

- List your loaded modules:

  `module list`

- To clear-out your added modules:

  `module purge` or **`module clear`**

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

The PBS Queue System

PBS Pro 11.1 is the resource management service used on the Palmetto Cluster

- Enables you to make more efficient use of your time through scripting computational tasks
- PBS takes care of running these tasks and returning the results
- If the cluster is full, PBS holds your tasks and runs them when the resources are available
- PBS ensures fair sharing of cluster resources (policy enforcement)
- PBS ensures optimal/efficient use of available resources

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

PBS Commands

Command(s)	Description
qsub *job-script*	Submit a batch job
qsub -I *resources*	Submit an interactive job
qstat -u *userid*	Check status of all of your jobs
qhold *jobID*	Put a job on hold (before it starts)
qrls *jobID*	Release a job from hold status
qdel *jobID*	Delete a job, running or not
qselect *criteria*	Select jobs by specified criteria

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

qstat –xf *[job ID]*

```
Job Id: 817100.pbs01
    Job_Name = ucaccc
    resources_used.cpupercent = 791
    resources_used.cput = 10:50:29
    resources_used.mem = 270832kb
    resources_used.ncpus = 32
    resources_used.vmem = 1170952kb
    resources_used.walltime = 01:21:34
    exec_host = node1026/0*8+node1027/0*8+node1028/0*8+node1029/0*8
    Error_Path = user001.palmetto.clemson.edu:/scratch/galen/amber/ucaccc/ucaccc.stderr
    Output_Path = user001.palmetto.clemson.edu:/scratch/galen/amber/ucaccc/ucaccc.stdout
    Resource_List.mem = 44gb
    Resource_List.mpiprocs = 32
    Resource_List.ncpus = 32
    Resource_List.nodect = 4
    Resource_List.place = free:shared
    Resource_List.select = 4:ncpus=8:mpiprocs=8:mem=11gb
    Resource_List.walltime = 72:00:00
    jobdir = /scratch/galen
    Variable_List = PBS_O_SYSTEM=Linux,PBS_O_SHELL=/bin/bash,
        PBS_O_HOME=/home/galen,PBS_O_HOST=user001.palmetto.clemson.edu,
        PBS_O_LOGNAME=galen,PBS_O_WORKDIR=/scratch/galen/amber/ucaccc,
        PBS_O_LANG=en_US.UTF-8,
        PBS_O_PATH=/scratch/galen/protg/mmtsb.v.Jul-31-2009/perl:/scratch/gale
        n/protg/mmtsb.v.Jul-31-2009/bin:/usr/lib64/qt-3.3/bin:/opt/pbs/default/
        bin:/opt/gold/bin:/opt/condor/bin:/opt/condor/sbin:/usr/local/bin:/bin:
        /usr/bin:/usr/local/sbin:/usr/sbin:/sbin:/opt/mx/bin:/home/galen/bin,
        PBS_O_MAIL=/var/spool/mail/galen,PBS_O_QUEUE=workq
```

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Some qsub Options

qsub -N name	Job name
qsub -q workq	Queue to assign job to
qsub -V	Export environment variables
qsub -v var=value	Expand upon environment variables
qsub -l (see below)	Resource list (hardware required)

chip_manufacturer=amd or **chip_manufacturer=intel**

chip_model=opteron or **chip_model=xeon**

chip_type=e5345 **chip_type=e5410** (l5420, x7542, 2356, 6172, etc.)

node_manufacturer=dell (hp, ibm, sun, dell)

For all available qsub options, see PBS User Guide, p. 75

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Example Batch Job Script File

```
#!/bin/bash
#PBS -N jobname
#PBS -l select=4:ncpus=8:mpiprocs=8
#PBS -l mem=11gb
#PBS -l chip_manufacturer=amd
#PBS -l walltime=2:00:00
#PBS -o stdout.txt
#PBS -e stderr.txt
#PBS -q workq
#PBS -M userid@clemson.edu

my.program.exe [arguments for my program]
```

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Notification Parameters

`#PBS -M [e-mail address]`

e-mail address can be a list of email addresses separated by commas

`#PBS -m bean` or **`-m be`** or **`-m e`**

Mail options: send an email when the job **b**egins, **e**nds, is **a**borted, or **n**o notification

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Handling Output Files

- The 'standard output' and 'error output' are sent to files named ***jobname.o987349*** and ***jobname.e987349*** (in your $PBS_O_WORKDIR directory).

- The following parameters can modify this behavior:

`#PBS -e [path/filename]`	error output file
`#PBS -o [path/filename]`	standard o/p
`#PBS -j eo`	merge the error and standard output
`#PBS -k eo`	keep error and standard output separate

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Using an Interactive Job

- Useful for debugging applications, short tests, or for computational steering

```
qsub -I

qsub -I -l walltime=2:00:00

qsub -I -l select=1:ncpus=8:mem=11gb

qsub -I -l walltime=2:00:00,select=2:npcus=8

qsub -I -l select=2:ncpus=8:mpiprocs=8:mem=15gb,walltime=2:00:00

qsub -I -q bigmem -l select=1:ncpus=1:mem=64gb,walltime=4:00:00
```

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

checkqueuecfg

```
Default (workq) routing queue configs


MaxRun/MaxQue            0-2 hrs(quick)  2-24 hrs(short) 24-72 hrs(long)
================================================================================
       1 nodes(single)   450/unlimited   450/unlimited   450/unlimited
    2-16 nodes(tiny)     30/unlimited    30/unlimited    30/unlimited
   17-64 nodes(small)    15/unlimited    15/unlimited    15/unlimited
  65-255 nodes(medium)   4/unlimited     4/unlimited     4/unlimited
256-1622 nodes(large)    1/unlimited     1/unlimited     1/unlimited
```

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Killing Jobs

One, or a few jobs:

```
qdel [jobID] [jobID] [jobID] ...
```

Kill all of your jobs:

```
qselect -u $USER | xargs qdel
```

Kill all of your queued jobs:

```
qselect -u $USER -s Q | xargs qdel
```

Kill all of your running jobs:

```
qselect -u $USER -s R | xargs qdel
```

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Palmetto User Support

- “Help” requests or technical questions can be submitted to the Palmetto Admin Staff and anyone who can help will respond.
- Computational scientists are available for consultation and training.
- We can help with compiling code and installing programs.
- We can also help with developing proposals that make use of Palmetto resources.

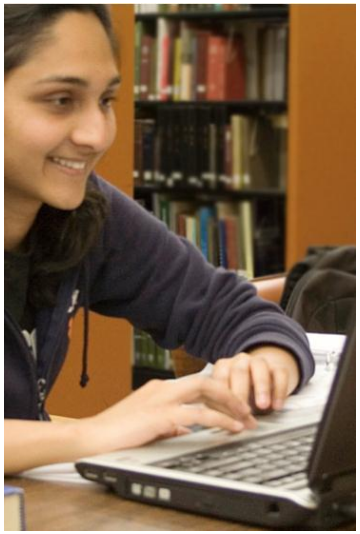

CLEMSON
COMPUTING AND
INFORMATION TECHNOLOGY

Final Notes

Palmetto Cluster User's Guide

PBS Professional 11.1 User Guide

- Do not run jobs on the head/login node (user001)
- Work in /scratch, then move all files when finished.
- When you log-in, read the MOTD
- Need help?

  E-mail ithelp@clemson.edu

  with the word "Palmetto" in the subject line…

  job ID or copy of your PBS job script helps a lot